

Diabetes Prediction using Machine Learning

---Maitreya Sameer Ganu , IISER Thiruvananthapuram

```
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd

data =
pd.read_csv('/kaggle/input/pima-indians-diabetes-database/diabetes.csv')
data.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Let's visualize our data

```
data
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1

4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

data.describe()

	Pregnancies	Glucose	BloodPressure	SkinThickness
Insulin \				
count	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458
std	3.369578	31.972618	19.355807	15.952218
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000
75%	6.000000	140.250000	80.000000	32.000000
max	17.000000	199.000000	122.000000	99.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64
8	Outcome	768 non-null	int64

```
dtypes: float64(2), int64(7)
```


```
memory usage: 54.1 KB
```

Clearly , there are no missing values and values which are null

```
plt.figure(figsize = (12,6))
```

```
sns.countplot(x = 'Outcome' , data = data)
```

```
<Axes: xlabel='Outcome', ylabel='count'>
```

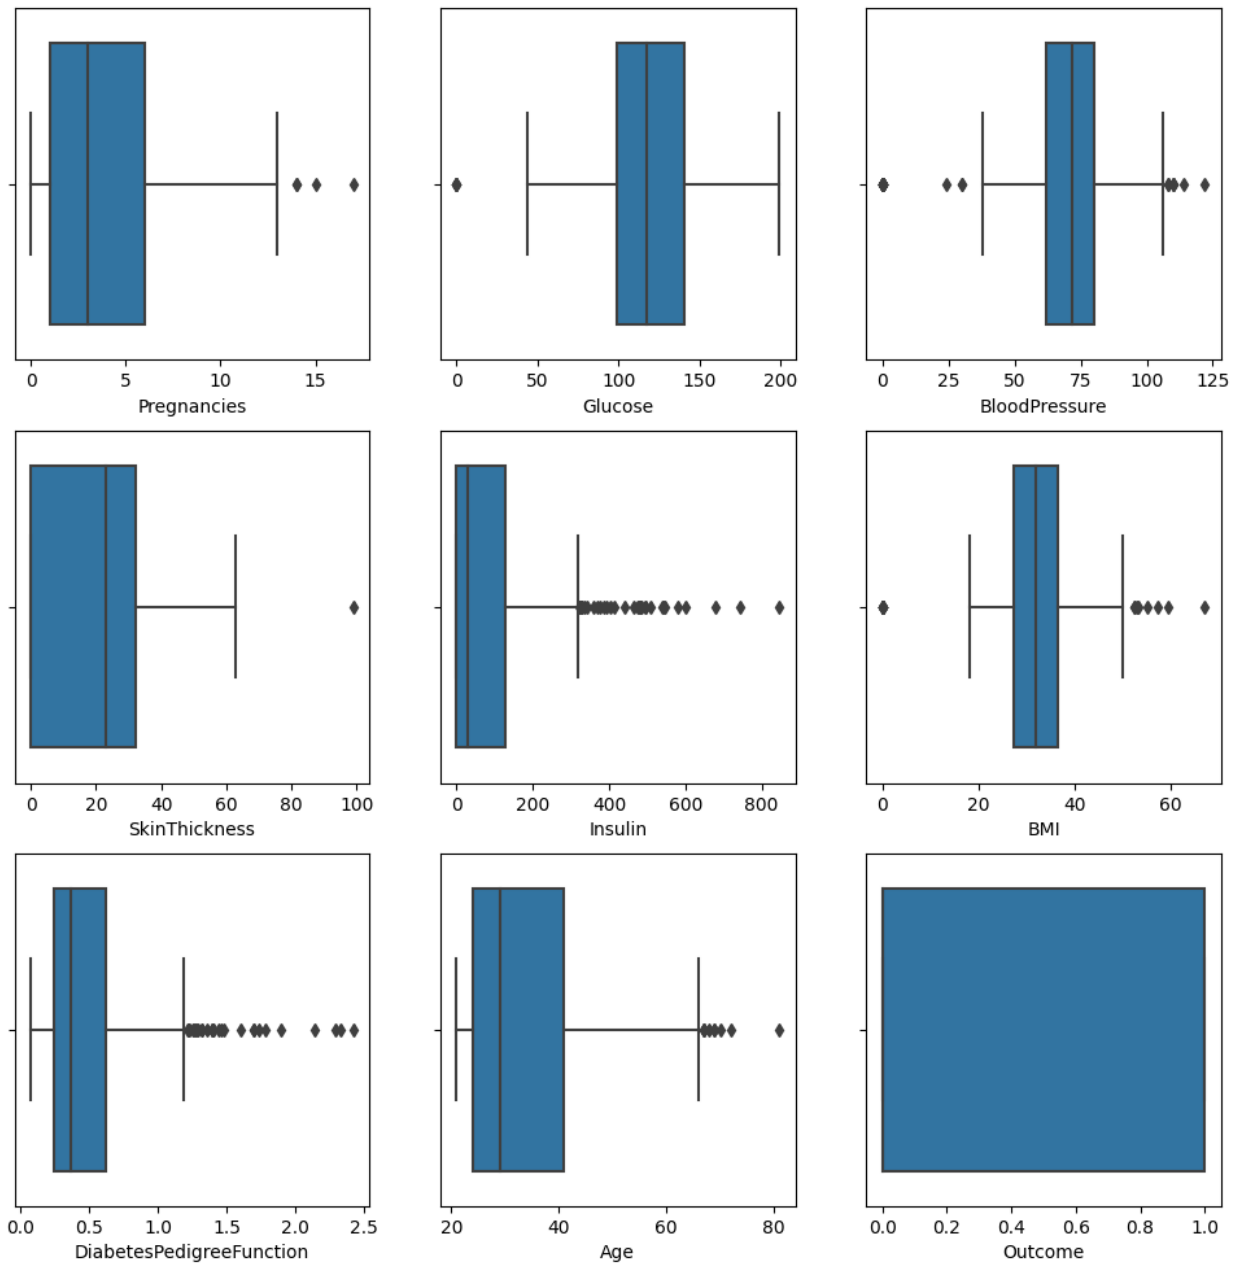


Let's observe any outliers

```
data.columns
```

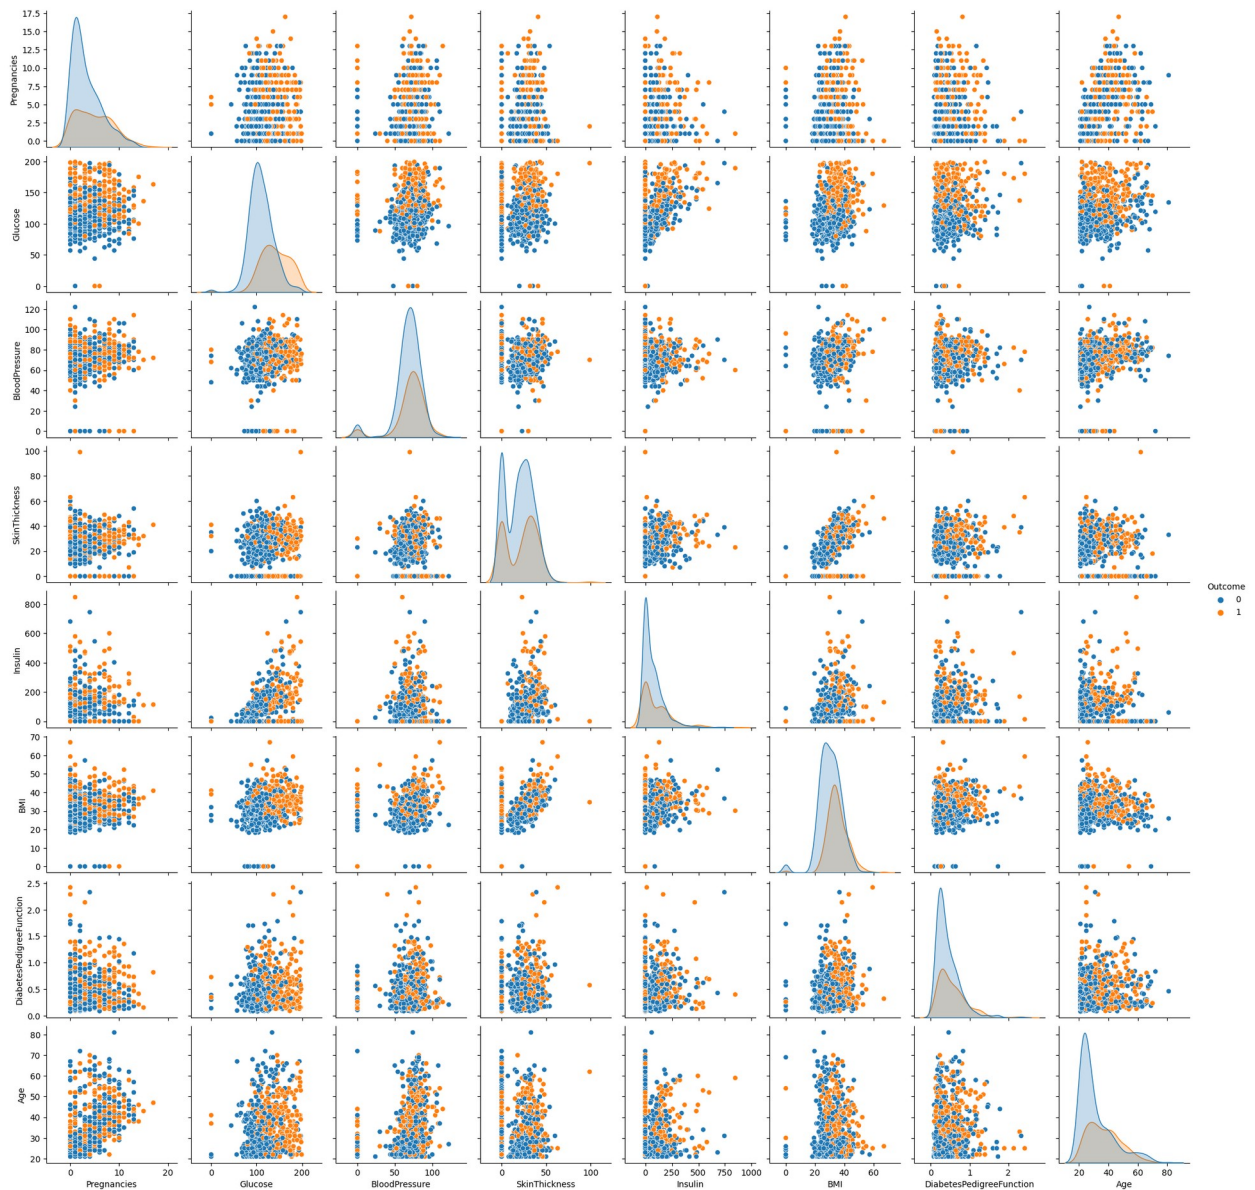
```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',  
      'Insulin',  
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
      dtype='object')
```

```
plt.figure(figsize = (12,12))  
for i,col in enumerate(['Pregnancies', 'Glucose', 'BloodPressure',  
                        'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age',  
                        'Outcome']):  
    plt.subplot(3,3 , i+1)  
    sns.boxplot(x = col , data = data)  
plt.show()
```

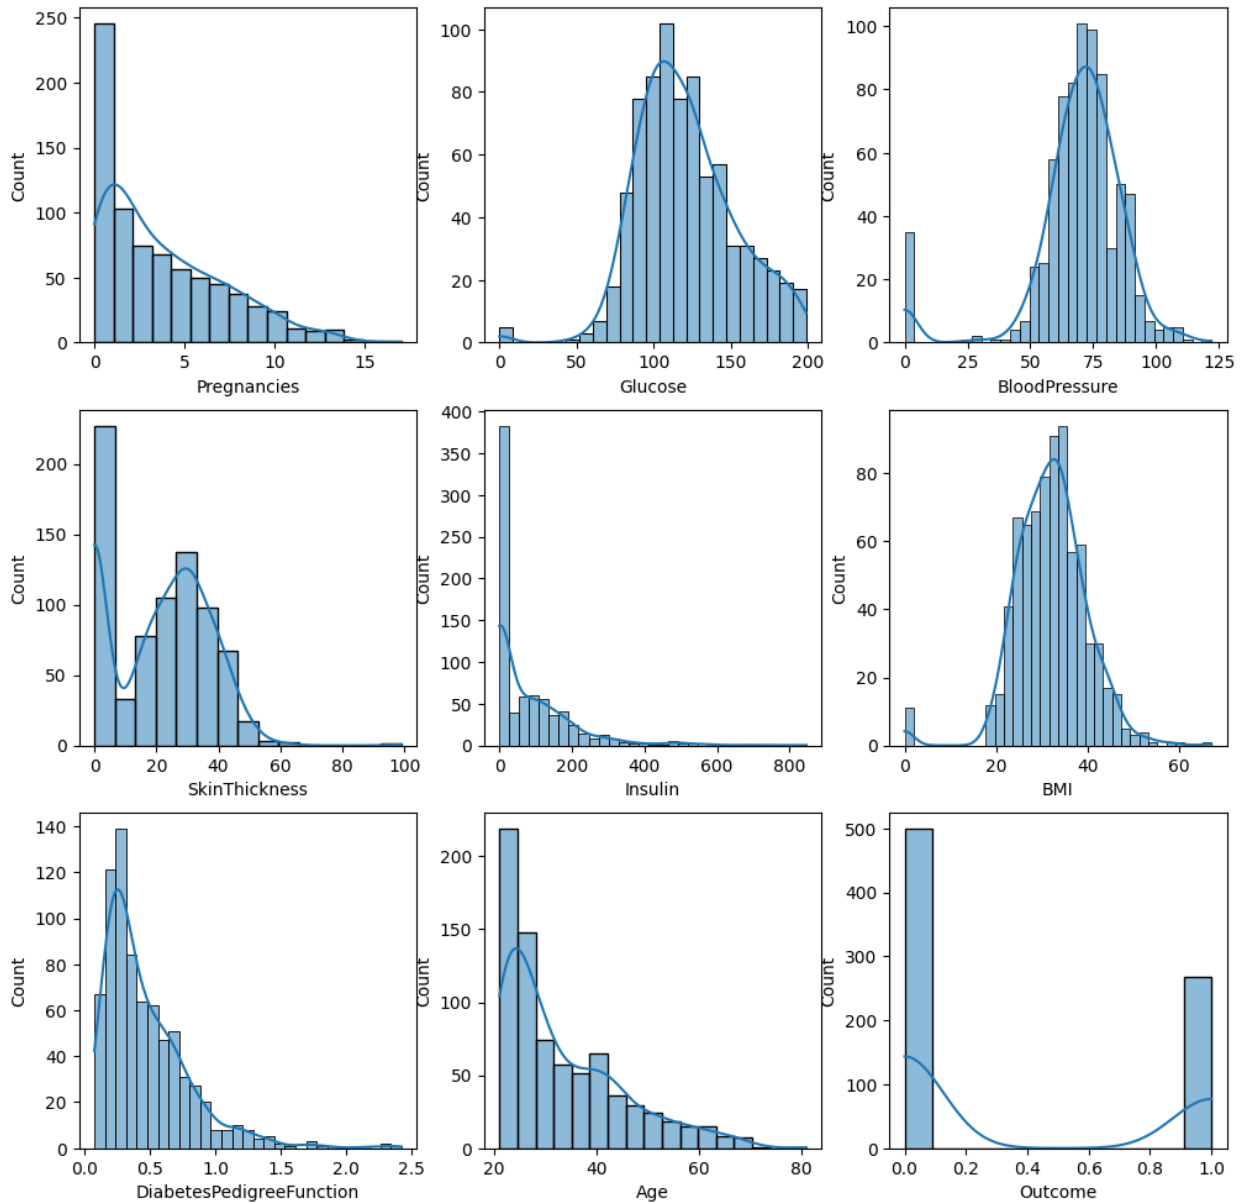


```
import warnings


warnings.filterwarnings('ignore')
sns.pairplot(data , hue = 'Outcome')
plt.show()
```



```
plt.figure(figsize = (12,12))
for i,col in enumerate(['Pregnancies', 'Glucose', 'BloodPressure',
                        'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age',
                        'Outcome']):
    plt.subplot(3,3 , i+1)
    sns.histplot(x = col , data = data , kde = True)
plt.show()
```



```
plt.figure(figsize = (12,12))
sns.heatmap(data.corr() , vmin = -1 , center = 0 , cmap = 'coolwarm' ,
annot = True)
plt.show()
```



Standard Scaling and Label Encodings

```
from sklearn.preprocessing import StandardScaler
import warnings

warnings.filterwarnings('ignore')
sc_X = StandardScaler()
X = pd.DataFrame(sc_X.fit_transform(data.drop(['Outcome'], axis =
1)), columns = ['Pregnancies', 'Glucose', 'BloodPressure',
'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age'])
```



```
X.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	0.639947	0.848324	0.149641	0.907270	-0.692891	0.204013
1	-0.844885	-1.123396	-0.160546	0.530902	-0.692891	0.684422
2	1.233880	1.943724	-0.263941	-1.288212	-0.692891	1.103255
3	-0.844885	-0.998208	-0.160546	0.154533	0.123302	0.494043
4	-1.141852	0.504055	-1.504687	0.907270	0.765836	1.409746

	DiabetesPedigreeFunction	Age
0	0.468492	1.425995
1	-0.365061	-0.190672
2	0.604397	-0.105584
3	-0.920763	-1.041549
4	5.484909	-0.020496

```
y = data['Outcome']
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2 ,
random_state = 0)
```

```
from sklearn.neighbors import KNeighborsClassifier
test_scores = []
train_scores = []
for i in range(1,15):
    knn = KNeighborsClassifier(i)
    knn.fit(X_train , y_train)
    train_scores.append(knn.score(X_train,y_train))
    test_scores.append(knn.score(X_test,y_test))
```

```
max_train_score = max(train_scores)
train_scores_index = [i for i,v in enumerate(train_scores) if v ==
max_train_score]
print("Maximum Train Score {} % and k = {}".format(max_train_score*100
, list(map(lambda x: x+1 , train_scores_index))))
```

```
Maximum Train Score 100.0 % and k = [1]
```

```
max_test_score = max(test_scores)
test_scores_index = [i for i,v in enumerate(test_scores) if v ==
max_test_score]
print("Maximum Test Score {} % and k = {}".format(max_test_score*100 ,
list(map(lambda x: x+1 , test_scores_index))))
```

Maximum Test Score 80.51948051948052 % and k = [5]

```
import warnings

warnings.filterwarnings('ignore')
plt.figure(figsize=(12, 5))
sns.lineplot(x=range(1, 15), y=train_scores, marker='o', label='Train scores')
sns.lineplot(x=range(1, 15), y=test_scores, marker='o', label='Test scores')
plt.title('Train vs Test Scores', fontsize=14)
plt.xlabel('Model Complexity', fontsize=12)
plt.ylabel('Score', fontsize=12)

plt.grid(True)
plt.legend()
plt.show()
```



Here , for k = 1 we are getting the highest train score and for k = 5 ,we are getting the highest test score

```
knn = KNeighborsClassifier(5)
knn.fit(X_train,y_train)
knn.score(X_test,y_test)

0.8051948051948052

from sklearn.metrics import confusion_matrix,classification_report
y_pred = knn.predict(X_test)
print(confusion_matrix(y_test,y_pred))
print(classification_report(y_test,y_pred))
```

```
[[94 13]
 [17 30]]
```

	precision	recall	f1-score	support
0	0.85	0.88	0.86	107
1	0.70	0.64	0.67	47
accuracy			0.81	154
macro avg	0.77	0.76	0.76	154
weighted avg	0.80	0.81	0.80	154

Clearly , we got a pretty accurate output.

Here , 20% of the data was used for testing purposes taking into consideration the Pareto's Principle.